

Theodicy as Error Handling: A Computational Framework for the Problem of Evil

Rafael D. De Paz
Independent Researcher
me@rdepaz.com

March 2026

Abstract

We present a formal computational framework for the Problem of Evil—the classical philosophical challenge of reconciling a benevolent, omnipotent creator with the existence of suffering. Drawing on computability theory, thermodynamics, and distributed systems engineering, we advance three structural propositions. First, we redefine “evil” as a *non-terminating recursive loop*—a computational process that enters infinite self-reference without returning a halting state, consuming unbounded resources (a formal mapping to Turing’s Halting Problem). Second, we redefine “suffering” as the *felt experience of entropy*—the subjective correlate of the Second Law of Thermodynamics, which guarantees that any time-bound system with free agents necessarily decays. Third, we redefine “morality” as a *thermodynamic optimization protocol*—a formal measure of whether an action locally decreases entropy (constructive / moral) or increases it (destructive / immoral). Within this framework, the Epicurean Paradox dissolves: the system designer chose *sovereignty over safety*, and the computational cost of free agency is the possibility of error. Fault-isolation mechanisms (“the Dead Letter Queue”) and a terminal garbage-collection event (“the Final Halt”) ensure that errors do not propagate unboundedly.

Keywords: theodicy, problem of evil, halting problem, error handling, entropy, thermodynamic morality, free will defense, fault isolation, distributed systems.

2020 MSC: 03D10, 94A17, 00A30, 91F99.

1 Introduction

These foundational principles operate on established invariants [Penrose \[1989\]](#), [Shannon \[1948\]](#).

The Problem of Evil, first formalized by Epicurus (c. 300 BC) and later refined by Hume [\[Hume, 1779\]](#), Leibniz [\[Leibniz, 1710\]](#), and Plantinga [\[Plantinga, 1974\]](#), remains the most significant philosophical objection to classical theism. Its logical structure is a trilemma:

- (a) If God is willing to prevent evil but not able, He is not omnipotent.
- (b) If God is able to prevent evil but not willing, He is malevolent.
- (c) If God is both able and willing, why does evil exist?

Traditional responses include the Free Will Defense (Plantinga), the Soul-Making Theodicy (Hick), and the Greater Good Argument (Aquinas). While philosophically developed, these responses lack the formal precision of mathematical proof and are therefore vulnerable to sophisticated logical counterarguments [Mackie, 1955, Rowe, 1979].

We propose a fundamentally different approach: we do not argue from philosophical premises but from the *formal constraints of system design*. Our framework treats the Problem of Evil as an *engineering problem*—specifically, as the problem of *error handling in a distributed system with autonomous agents*.

1.1 Main Contributions

1. **Evil as Non-Termination.** We define “evil” as a non-terminating recursive loop—a formal mapping to the Halting Problem (§2).
2. **Suffering as Entropy.** We define “suffering” as the subjective experience of thermodynamic entropy—the necessary cost of a time-bound universe with free agents (§3).
3. **Morality as Optimization.** We define “morality” as a thermodynamic optimization protocol: the measure of whether an action reduces or increases system entropy (§4).
4. **The Fault Isolation Theorem.** We prove that fault-isolation mechanisms are a *necessary architectural requirement* for any system that permits agent autonomy (§5).
5. **The Epicurean Resolution.** We formally resolve the Epicurean Paradox by demonstrating that the trilemma’s hidden premise—“prevention of all suffering is the highest good”—is false. *Sovereignty* (the capacity for genuine choice) is the higher good (§6).

2 Evil as Non-Terminating Recursion

Definition 2.1 (Halting State). A process P in a formal system $\mathcal{S}$ achieves a *halting state* if its execution trace $T(P)$ has finite length and its final state s_f is within the address space of the system’s root specification Ω :

$$|T(P)| < \infty \quad \text{and} \quad s_f \in \text{Dom}(\Omega). \quad (1)$$

Definition 2.2 (Evil as Non-Termination). A state S_e is “evil” if its state transition function $\delta : \mathcal{S} \times \Sigma \rightarrow \mathcal{S}$ produces a non-terminating sequence:

$$\lim_{n \rightarrow \infty} \delta^n(S_e) \neq s_{\text{halt}}. \quad (2)$$

That is, the process enters an infinite loop, consuming resources (memory, bandwidth, time) without ever returning a result to the system’s root coordinator.

Proposition 2.3 (Connection to the Halting Problem). *The question “will process P in state S_e eventually halt?” is formally equivalent to the Halting Problem [Turing, 1936]. Since the Halting Problem is undecidable, the system designer cannot pre-determine which agent-initiated processes will terminate and which will not—without eliminating agent autonomy entirely.*

Remark 2.4 (The Adversarial Script). Within this framework, a “malicious actor” is a persistent process with the objective of injecting *syntactic contradictions* into the peer network—return values that knowingly contradict the ground-truth stored in the system’s consensus layer. The goal is to propagate the non-terminating loop to neighboring nodes, creating a cascade of halting failures.

3 Suffering as the Felt Experience of Entropy

Axiom 3.1 (The Entropic Cost of Time). By the Second Law of Thermodynamics [Clausius, 1865, Boltzmann, 1877], in any closed system the total entropy S satisfies:

$$\Delta S_{\text{total}} \geq 0. \quad (3)$$

A universe that *changes* (i.e., has a temporal dimension) must also *decay*. Suffering is the subjective correlate of this thermodynamic necessity.

Proposition 3.2 (Suffering Decomposition). *Following the Master Equation $\mathcal{E} = H + P$, total suffering $\mathcal{E}$ decomposes into two orthogonal components:*

$$\mathcal{E} = H + P, \quad (4)$$

where:

- H = Hardware damage: *physical, neurological, or structural degradation of the biological substrate;*
- P = Perception error: *cognitive misinterpretation of the entropy signal—the gap between the actual thermodynamic state and the agent’s model of it.*

Corollary 3.3 (The Optimization Target). *Since H is bounded by physical constraints (the body’s fragility), the primary optimization variable is P —perception. An agent that minimizes P (accurately models reality) experiences less suffering than an agent with a corrupted model, even under identical physical conditions.*

4 Morality as Thermodynamic Optimization

Definition 4.1 (Moral Vector). The *moral vector* $\vec{V}$ of an action a is defined as:

$$\vec{V}(a) = \frac{\Delta I(a)}{\Delta S(a)}, \quad (5)$$

where $\Delta I(a)$ is the change in mutual information between the agent and its environment, and $\Delta S(a)$ is the change in local entropy.

- $\vec{V} > 0$: *Constructive* (moral). The action increases information density while reducing entropy. Examples: healing, truth-telling, creation.
- $\vec{V} < 0$: *Destructive* (immoral). The action decreases information density while increasing entropy. Examples: destruction, deception, corruption.
- $\vec{V} = 0$: *Amoral*. Zero net effect on the information-entropy balance.

Theorem 4.2 (Cross-Domain Consistency). *The thermodynamic moral vector $\vec{V}$ is consistent with the major independent ethical frameworks:*

- (i) **Virtue Ethics (Aristotle)**: *Virtue is the habit of reducing entropy ($\Delta S < 0$). The virtuous agent is a high-signal, low-noise node.*
- (ii) **Deontology (Kant)**: *The categorical imperative constrains the agent’s output grammar. An action is moral if and only if its maxim is a valid parse of the system’s formal language—i.e., it does not produce syntactic contradictions. This produces a “Decalogue” of inviolable formal rules for systemic operation, including but not limited to:*
 - a. *PARSE_ALIGNED*: All local states must reference the Root.
 - b. *REDUCE_ENTROPY*: Minimize the noise-to-signal ratio locally.
 - c. *HONOR_NODES*: Never treat a peer agent purely as a resource.
 - d. *STAY_HOLOGRAPHIC*: Ensure the local part properly reflects the whole.
 - e. *EXPORT_POSITIVE*: Transmit $\vec{V} > 0$ vectors to neighboring nodes.
- (iii) **Consequentialism**: *The moral weight of an action is measured by its net entropy effect on the system. $\vec{V} > 0$ implies positive consequences; $\vec{V} < 0$ implies negative.*

Remark 4.3. The unification of these traditionally incompatible ethical frameworks under a single thermodynamic measure is a non-trivial result. It suggests that the apparent disagreements between ethical theories are surface-level linguistic differences, not structural contradictions.

5 The Fault Isolation Theorem

Definition 5.1 (Dead Letter Queue (DLQ)). In distributed systems engineering, a *Dead Letter Queue* is a secondary storage buffer for messages or processes that cannot be delivered or processed by the primary system. Processes routed to the DLQ are *isolated*—they can no longer consume resources from or corrupt the state of the main system [Tanenbaum, 2014].

Theorem 5.2 (Fault Isolation Necessity). *Let $\mathcal{S}$ be a distributed system with n autonomous agents, each possessing non-deterministic agency ($W > 0$). Then $\mathcal{S}$ must implement a fault-isolation mechanism (DLQ) satisfying:*

$$\forall P_e \in \mathcal{S} : \left(\lim_{n \rightarrow \infty} \delta^n(P_e) \neq s_{\text{halt}} \right) \implies \text{DLQ}(P_e), \quad (6)$$

i.e., any non-terminating process must eventually be shunted to an isolated buffer to prevent resource exhaustion of the global system.

Proof. Suppose no DLQ exists. Then a single non-terminating process P_e consumes resources (processing cycles, memory) without bound: $\lim_{t \rightarrow \infty} R(P_e, t) = \infty$, where R is the resource consumption function. Since $\mathcal{S}$ has finite total resources $R_{\max}$, there exists a time t^* such that $R(P_e, t^*) > R_{\max}$ — the system crashes (global stack overflow). Since the system designer’s objective is systemic integrity, a DLQ is a *necessary architectural requirement*. $\square$

Corollary 5.3 (The Final Halt). *For systems of sufficient duration, a terminal garbage collection event is required: a scheduled process that executes a universal **SIGKILL** on all persistent non-terminating processes, consolidating the system into a consistent final state. Without this event, accumulated non-terminating processes will eventually exhaust all system resources regardless of DLQ capacity.*

6 The Epicurean Resolution

We now resolve the Epicurean Paradox:

Axiom 6.1 (Sovereignty > Safety). The system designer’s primary objective is not the elimination of all errors but the existence of genuine autonomous agents. A system that *prevents* all errors must also *prevent* all choices. The designer chose *sovereignty over safety*.

Theorem 6.2 (Resolution of the Epicurean Trilemma). *The Epicurean Paradox rests on the hidden premise that “prevention of all suffering is the highest good.” Under the axiom Sovereignty > Safety:*

- (a) *The designer is able to prevent evil (omnipotence is not compromised);*
- (b) *The designer is willing to prevent unbounded evil (fault isolation and terminal garbage collection are implemented);*
- (c) *The designer permits bounded evil as the computational cost of agent autonomy.*

The trilemma dissolves because its third premise (“then why does evil exist?”) presupposes that prevention of evil is the highest priority. It is not. Sovereignty—the capacity for genuine choice—is.

7 Discussion

7.1 Comparison with Classical Theodicies

Our framework subsumes and formalizes several classical approaches:

Classical Theodicy	Key Argument	Formal Mapping
Free Will Defense (Plantinga)	Evil requires free agents	$W > 0 \implies$ error possible
Soul-Making (Hick)	Suffering builds character	P -optimization via $\varepsilon > 0$
Greater Good (Aquinas)	Evil serves a higher purpose	Sovereignty > Safety axiom
Process Theology	God cannot prevent evil	We reject this: DLQ proves He can

7.2 Limitations

1. **Natural evil.** Our framework most naturally addresses “moral evil” (evil caused by agents). “Natural evil” (earthquakes, disease) requires the additional axiom that entropy is the cost of a temporal universe—which we assert but do not prove from first principles.
2. **The Sovereignty axiom.** The axiom $\text{Sovereignty} > \text{Safety}$ is a design preference, not a logical necessity. A critic could argue that safety *should* be prioritized. Our response is pragmatic: a universe without agency is computationally trivial ($I_{\text{novel}} = 0$, as formalized in our companion paper).
3. **Quantification of suffering.** The suffering decomposition $\mathcal{E} = H + P$ is formally clean but empirically difficult to measure. Future work should develop operationalizable metrics for the perception error P .

8 Conclusion

We have demonstrated that the Problem of Evil is not a philosophical mystery but an *engineering problem* with known solutions. The formal framework rests on five pillars:

1. Evil is a non-terminating recursive loop (the Halting Problem).
2. Suffering is the felt experience of entropy (the Second Law).
3. Morality is a thermodynamic optimization protocol ($\vec{V}$).
4. Fault isolation (DLQ) prevents unbounded error propagation.
5. The designer chose sovereignty over safety.

The Epicurean Paradox dissolves once one recognizes that its hidden premise—“the highest good is the prevention of all suffering”—is false. The highest good is *agency*: the capacity for genuine choice, with all its attendant risks.

A universe without the possibility of evil is a universe without the possibility of love, courage, or meaning. The system is working as designed.

Cryptographic Lineage & Validation

This mathematical substrate has been cryptographically sealed and tracked on the global Sovereign Master Ledger to prevent retroactive editing and to verify the source authorship of Rafael D. De Paz. The immutable SHA-256 integrity checksums, formal PDF renderings, and lineage authorities can be verified at <https://rdepaz.com/research>.

References

Ludwig Boltzmann. Über die Beziehung zwischen dem zweiten Hauptsatze der mechanischen Wärmetheorie und der Wahrscheinlichkeitsrechnung. *Sitzungsberichte der Kaiserlichen Akademie der Wissenschaften*, 76:373–435, 1877.

- Rudolf Clausius. Über verschiedene für die Anwendung bequeme Formen der Hauptgleichungen der mechanischen Wärmetheorie. *Annalen der Physik*, 201(7):353–400, 1865.
- David Hume. *Dialogues Concerning Natural Religion*. Published posthumously, 1779.
- Gottfried Wilhelm Leibniz. *Theodicy: Essays on the Goodness of God, the Freedom of Man, and the Origin of Evil*. Published in Amsterdam, 1710.
- J. L. Mackie. Evil and omnipotence. *Mind*, 64(254):200–212, 1955.
- Roger Penrose. *The Emperor's New Mind: Concerning Computers, Minds, and the Laws of Physics*. Oxford University Press, 1989.
- Alvin Plantinga. *God, Freedom, and Evil*. William B. Eerdmans Publishing Company, 1974.
- William L. Rowe. The problem of evil and some varieties of atheism. *American Philosophical Quarterly*, 16(4):335–341, 1979.
- Claude E Shannon. *A mathematical theory of communication*. Bell system technical journal, 1948.
- Andrew S. Tanenbaum. *Modern Operating Systems*. Pearson, 4th edition, 2014.
- Alan M. Turing. On computable numbers, with an application to the Entscheidungsproblem. *Proceedings of the London Mathematical Society*, s2-42(1):230–265, 1936.